

ECE408 Project Report

Team: sthBigBig

Members: Guanchen He(ghe10), Guxin Jin(gjin7), Yingyi Zhang(yingyiz2)

1.3 NVPROF Profile

```
—307— NVPROF is profiling process 307, command: python /src/m1.2.py
Loading model...[05:51:06] src/operator/././cudnn_algorreg-inl.h:112: Running performance tests to find the best convolution algorithm, this can take a while... (setting env variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done
EvalMetric: {'accuracy': 0.8673}
—307— Profiling application: python /src/m1.2.py
—307— Profiling result:
Time(%)   Time      Calls      Avg      Min      Max      Name
36.46%  49.296ms      1  49.296ms  49.296ms  49.296ms void cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, float const *, int, cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>*, float const *, kernel_conv_params, int, float, float, int, float const *, float const *, int, int)
28.20%  38.131ms      1  38.131ms  38.131ms  38.131ms sgemm_sm35_ldg_tn.128x8x256x16x32
14.33%  19.384ms      2  9.6920ms  455.22us  18.929ms void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
10.64%  14.392ms      1  14.392ms  14.392ms  14.392ms void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
5.70%   7.7031ms     13  592.55us  1.6000us  5.6581ms [CUDA memcopy HtoD]
2.68%   3.6252ms      1  3.6252ms  3.6252ms  3.6252ms sgemm_sm35_ldg_tn.64x16x128x8x32
0.81%   1.0991ms      1  1.0991ms  1.0991ms  1.0991ms void mshadow::cuda::SoftmaxKernel<int=8, float, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>(mshadow::gpu, int=2, unsigned int)
0.55%   738.45us     12  61.537us  2.0480us  372.57us void mshadow::cuda::MapPlanKernel-mshadow::sv::saveto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::ScalarExp<float>, float>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.32%   430.55us      2  215.28us  17.119us  413.43us void mshadow::cuda::MapPlanKernel-mshadow::sv::plusto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::Broadcast1DExp-mshadow::Tensor-mshadow::gpu, int=1, float>, float, int=2, int=1>, float>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.29%   391.51us      1  391.51us  391.51us  391.51us sgemm_sm35_ldg_tn.32x16x64x8x16
0.02%   22.911us      1  22.911us  22.911us  22.911us void mshadow::cuda::MapPlanKernel-mshadow::sv::saveto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::ReduceWithAxisExp-mshadow::red::maximum, mshadow::Tensor-mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01%    9.6000us      1  9.6000us  9.6000us  9.6000us [CUDA memcopy DtoH]
```

Figure 1 screenshot of profile application

Table 1 several time-consuming kernels

Time(%)	Time	Name
36.46%	49.296ms	implicit_convolve_sgemm
28.20%	38.131ms	sgemm_sm35_ldg_tn
14.33%	19.384ms	activation_fw_4d_kernel
10.64%	14.392ms	pooling_fw_4d_kernel
5.70%	7.7031ms	cuda memcopy HtoD
2.68%	3.6252ms	sgemm_sm35_ldg_tn

From the table we find that the forward activation and pooling part could be optimized by leveraging parallel algorithm or techniques discussed in this lecture.

```

--307-- API calls:
Time(%)   Time           Calls      Avg      Min      Max      Name
46.79%    1.98587s        18    110.33ms  16.632us  992.62ms  cudaStreamCreateWithFlags
29.09%    1.23443s        10    123.44ms  827ns    345.95ms  cudaFree
20.47%    868.93ms        24    36.205ms  243.29us  861.74ms  cudaMemGetInfo
2.99%     127.03ms        25    5.0813ms  5.2510us  82.540ms  cudaStreamSynchronize
0.37%     15.719ms        8     1.9648ms  12.855us  5.7855ms  cudaMemcpy2DAsync
0.16%     6.9237ms        42    164.85us  9.2990us  1.2400ms  cudaMalloc
0.03%     1.3636ms        4     340.89us  336.01us  348.48us  cuDeviceTotalMem
0.02%     882.37us       352    2.5060us  243ns    71.116us  cuDeviceGetAttribute
0.01%     597.14us       114    5.2380us  623ns    159.46us  cudaEventCreateWithFlags
0.01%     543.28us       23    23.620us  10.228us  104.52us  cudaLaunch
0.01%     476.00us        6    79.332us  57.361us  124.03us  cudaMemcpy
0.01%     440.51us        4    110.13us  56.651us  165.33us  cudaStreamCreate
0.01%     350.52us        2    175.26us  56.939us  293.58us  cudaStreamCreateWithPriority
0.00%     103.51us        4    25.876us  15.561us  32.709us  cuDeviceGetName
0.00%     84.422us       32    2.6380us  624ns    7.1390us  cudaSetDevice
0.00%     71.908us      110    653ns    411ns    2.2620us  cuDeviceGetAttribute
0.00%     63.964us      147    435ns    256ns    1.6370us  cudaSetupArgument
0.00%     25.443us       23    1.1060us  496ns    3.7500us  cudaConfigureCall
0.00%     16.948us       10    1.6940us  1.1310us  2.3760us  cudaGetDevice
0.00%     11.324us        1    11.324us  11.324us  11.324us  cudaBindTexture
0.00%     9.1350us       16    570ns    365ns    938ns    cudaPeekAtLastError
0.00%     5.5870us        1    5.5870us  5.5870us  5.5870us  cudaStreamGetPriority
0.00%     5.0210us        6    836ns    275ns    2.1010us  cuDeviceGetCount
0.00%     4.0940us        2    2.0470us  1.4600us  2.6340us  cudaEventRecord
0.00%     3.6890us        2    1.8440us  1.4830us  2.2060us  cudaStreamWaitEvent
0.00%     3.5460us        2    1.7730us  1.6860us  1.8600us  cudaDeviceGetStreamPriorityRange
0.00%     3.5240us        6    587ns    283ns    1.1230us  cuDeviceGet
0.00%     3.3930us        3    1.1310us  1.0670us  1.1640us  cuInit
0.00%     3.0760us        6    512ns    289ns    760ns    cudaGetLastError
0.00%     2.1040us        3    701ns    650ns    744ns    cuDriverGetVersion
0.00%     1.5270us        1    1.5270us  1.5270us  1.5270us  cudaUnbindTexture
0.00%     1.2840us        1    1.2840us  1.2840us  1.2840us  cudaGetDeviceCount
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-8c359489-3b52-4933-8995-449377d0ac26.tar.gz. The data will be pre
sent for only a short duration of time.
* Server has ended your request.

```

Figure 2 screenshot of API calls.

Table 2 some time-consuming API calls

Time(%)	Time	Name
46.79%	1.9858s	cudaStreamCreateWithFlag
29.09%	1.2344s	cudaFree
20.47%	868.93ms	cudaMemGetInfo
2.99%	127.03ms	cudaStreamSynchronize

2.1 Simple CPU implementation

In this step, we implemented a CPU convolutional kernel. We flowed the forward convolution described in Chapter 16 of the textbook. The implementation is a for-loop that loop through all the computation positions and produces the convolution result. The classification results with our convolution kernel are presented in the following figures.

```

--0.11.0)
Installing collected packages: mxnet
  Running setup.py develop for mxnet
Successfully installed mxnet
* Running python /src/m2.1.py
Loading fashion-mnist data... done
Loading model... done
Op Time: 18.190609
Correctness: 0.8562 Model: ece408-high
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-3dbe502a-67f9-4044-ad5e-3a2d78393a0c.tar.gz. The data will be present for only a short duration of time.
* Server has ended your request.

```

Figure 3 ece408-high model execution time and accuracy

```

Installing collected packages: mxnet
  Running setup.py develop for mxnet
Successfully installed mxnet
* Running python /src/m2.1.py ece408-low 100
Loading fashion-mnist data... done
Loading model... done
Op Time: 0.201315
Correctness: 0.63 Model: ece408-low
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-7d417298-b51b-4a98-9005-7097a0b9ed36.tar.gz. The data will be present for only a short duration of time.
* Server has ended your request.

```

Figure 4 ece408-low model execution time and accuracy

3 Simple GPU implementation and oprimization plan

3.1 Simple GPU implementation

In this step, we implemented a simple GPU forward implementation. This implementation applies the convolution strategy discussed in ECE408 classes, i.e. tiled convolution. The implementation uses shared memory to apply convolution operation without extensive global memory loads. Figure 5 presents the performance of this version of implementation printed with nvprof. The kernel function tasks about 160ms to finish.

```

*Running nvprof python m3.1.py
Loading fashion-mnist data... done
==311== NVPROF is profiling process 311, command: python m3.1.py
Loading model... done
Op Time: 0.161134
Correctness: 0.8562 Model: ece408-high
==311== Profiling application: python m3.1.py
==311== Profiling result:
Time(%)    Time    Calls    Avg      Min      Max  Name
65.11%    161.03ms      1  161.03ms  161.03ms  161.03ms  void mxnet::op::forward_kernel<mshadow::gpu, float>
(float*, mxnet::op::forward_kernel<mshadow::gpu, float> const *, mxnet::op::forward_kernel<mshadow::gpu, floa
t> const , int, int, int, int, int, int)

```

Figure 5 performance of our simple GPU forward implementation.

During the forward pass of the convolutional neural network, the weights in the kernel won't change. A straightforward optimization strategy is putting the kernel into constant memory to reduce the time consumption of accessing global memory. Since the kernel weights are already in global memory, we need to use **cudaMemcpyToSymbol()** function with copy type **cudaMemcpyDeviceToDevice** to copy kernel weights into constant memory. The performance of tiled convolution with constant memory is presented in Figure 6.

```

*Running nvprof python m3.1.py
Loading fashion-mnist data... done
==311== NVPROF is profiling process 311, command: python m3.1.py
Loading model... done
Op Time: 0.147708
Correctness: 0.8562 Model: ece408-high
==311== Profiling application: python m3.1.py
==311== Profiling result:
Time(%)    Time    Calls    Avg      Min      Max  Name
63.23%    147.54ms      1  147.54ms  147.54ms  147.54ms  void mxnet::op::forward_kernel<mshadow::gpu, float>
(float*, mxnet::op::forward_kernel<mshadow::gpu, float> const *, mxnet::op::forward_kernel<mshadow::gpu, floa
t> const , int, int, int, int, int, int)

```

Figure 6 performance of tiled convolution with constant memory.

According to the result produced by nvprof, this simple optimization reduce the forward time from ~160ms to ~147ms. Further optimization plans are discussed in next section.

3.2 Optimization plan

Besides the basic optimization with constant memory, we also discussed about the further optimization strategies. Chapter 16 mainly presents two kinds of optimization techniques: reduce the convolution to matrix multiplication and use FFT for convolution. We print the size of our kernel in our experiment. Since the kernel size is only 5X5, FFT may not provide great performance improvement in our convolution task. The reduction strategy doesn't have this kind of restriction. Thus we decide to apply this strategy and reduce our convolution to matrix multiplication.

Reducing convolution to matrix multiplication requires two steps to produce the convolution result. First, the kernels and the input data batch are converted into two large matrix in unroll step. Then a matrix multiplication is conducted to produce the final result. Figure 7 exhibits the basic process of this strategy.

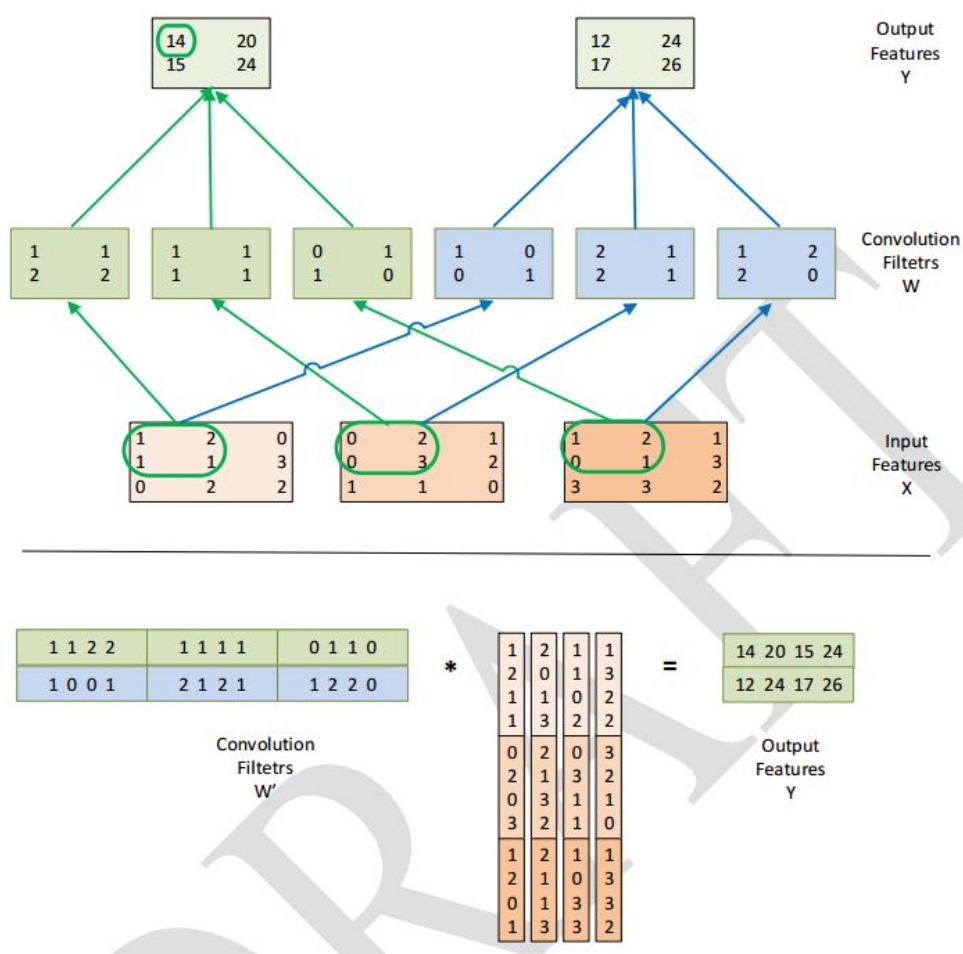


Figure 7. basic process of reducing convolution to matrix multiplication [1]

Our next optimization step is adding a new kernel to unroll the convolution kernels and the input data into two large matrices. The forward kernel will be responsible for the matrix multiplication. The matrix multiplication will be computed with tiled matrix multiplication discussed in our course.

We also discussed about some implementation details of the matrix multiplication strategy. According to our computation, the unrolled matrix can't be put into shared memory due to its large size. This might require multiple extra accesses on global memory. A possible

solution to this problem is combining the two steps in one kernel and create unrolled matrix on the fly and do the computation without writing the unrolled matrix to global memory. Besides, we plan to use streams to further speed up the computation in this combined kernel (since the computation of each unrolled sub-matrix doesn't rely on the results of others). In the multiplication part, we may use transpose to further utilize the burst potential. This part might be further optimized with non-square tile size. If the tile size is well designed, we may avoid divergence in the computation.

Note that this is just a plan of optimization. Some details might be changed during implementation.

Contribution: All the members thoroughly discussed the problem and distributed the task reasonably. The tasks are distributed as following

Milestone 1 : experiments done by Yingyi Zhang(yingyiz2), Guxin Jin(gjin7) and Guanchen He(ghe10).

Milestone 2: CPU code by Guanchen He(ghe10), experiment done by Yingyi Zhang(yingyiz2), report wrote by Guxin Jin(gjin7).

Milestone 3: GPU code by Yingyi Zhang(yingyiz2, optimization with constant memory done by Guanchen He(ghe10), report wrote by Guxin Jin(gjin7) and Guanchen He(ghe10).

The optimization plan was discussed and proposed by Yingyi Zhang(yingyiz2), Guxin Jin(gjin7) and Guanchen He(ghe10).

Reference

[1] Textbook Chapter 16 from course website. URL:<https://wiki.illinois.edu/wiki/display/ECE408Fall2017/Textbook+Chapters>