

ECE408 Project Report

Team: sthBigBig

Members: Guanchen He(ghe10), Guxin Jin(gjin7), Yingyi Zhang(yingyiz2)

1.3 NVPROF Profile

```
—307— NVPROF is profiling process 307, command: python /src/ml.2.py
Loading model...[05:51:06] src/operator/././cudnn_algoreg-inl.h:112: Running performance tests to find the best convolution algorithm, this can take a while... (setting env variable MXNET_CUDNN_AUTOTUNE_DEFAULT to 0 to disable)
done
EvalMetric: {'accuracy': 0.8673}
—307— Profiling application: python /src/ml.2.py
—307— Profiling result:
Time(%)   Time      Calls      Avg      Min      Max      Name
36.46% 49.296ms      1 49.296ms 49.296ms 49.296ms void cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=3, int=1, bool=1, bool=0, bool=1>(int, int, float const *, int, cudnn::detail::implicit_convolve_sgemm<float, int=1024, int=5, int=5, int=3, int=3, int=1, bool=1, bool=0, bool=1>*, float const *, kernel_conv_params, int, float, float, int, float const *, float const *, int, int)
28.20% 38.131ms      1 38.131ms 38.131ms 38.131ms sgemm_sm35_ldg_tn.128x8x256x16x32
14.33% 19.384ms      2 9.6920ms 455.22us 18.929ms void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)
10.64% 14.392ms      1 14.392ms 14.392ms 14.392ms void cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>(cudnnTensorStruct, float const *, cudnn::detail::pooling_fw_4d_kernel<float, float, cudnn::detail::maxpooling_func<float, cudnnNanPropagation_t=0>, int=0>, cudnnTensorStruct*, cudnnPoolingStruct, float, cudnnPoolingStruct, int, cudnn::reduced_divisor, float)
5.70% 7.7031ms     13 592.55us 1.6000us 5.6581ms [CUDA memcp HtoD]
2.68% 3.6252ms      1 3.6252ms 3.6252ms 3.6252ms sgemm_sm35_ldg_tn.64x16x128x8x32
0.81% 1.0991ms      1 1.0991ms 1.0991ms void mshadow::cuda::SoftmaxKernel<int=8, float, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>>(mshadow::gpu, int=2, unsigned int)
0.55% 738.45us     12 61.537us 2.0480us 372.57us void mshadow::cuda::MapPlanKernel-mshadow::sv::saveto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::ScalarExp<float>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.32% 430.55us      2 215.28us 17.119us 413.43us void mshadow::cuda::MapPlanKernel-mshadow::sv::plusto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::Broadcast1DExp-mshadow::Tensor-mshadow::gpu, int=1, float>, float, int=2, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.29% 391.51us      1 391.51us 391.51us 391.51us sgemm_sm35_ldg_tn.32x16x64x8x16
0.02% 22.911us      1 22.911us 22.911us 22.911us void mshadow::cuda::MapPlanKernel-mshadow::sv::saveto, int=8, mshadow::expr::Plan-mshadow::Tensor-mshadow::gpu, int=2, float>, float>, mshadow::expr::Plan-mshadow::expr::ReduceWithAxisExp-mshadow::red::maximum, mshadow::Tensor-mshadow::gpu, int=3, float>, float, int=3, bool=1, int=2>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=2)
0.01% 9.6000us      1 9.6000us 9.6000us 9.6000us [CUDA memcp DtoH]
```

Figure 1 screenshot of profile application

Table 1 several time-consuming kernels

Time(%)	Time	Name
36.46%	49.296ms	implicit_convolve_sgemm
28.20%	38.131ms	sgemm_sm35_ldg_tn
14.33%	19.384ms	activation_fw_4d_kernel
10.64%	14.392ms	pooling_fw_4d_kernel
5.70%	7.7031ms	cuda memcpy HtoD
2.68%	3.6252ms	sgemm_sm35_ldg_tn

From the table we find that the forward activation and pooling part could be optimized by leveraging parallel algorithm or techniques discussed in this lecture.

```

--307-- API calls:
Time(%)   Time           Calls      Avg      Min      Max      Name
46.79%    1.98587s        18    110.33ms  16.632us  992.62ms  cudaStreamCreateWithFlags
29.09%    1.23443s        10    123.44ms  827ns    345.95ms  cudaFree
20.47%    868.93ms        24    36.205ms  243.29us  861.74ms  cudaMemGetInfo
2.99%     127.03ms        25    5.0813ms  5.2510us  82.540ms  cudaStreamSynchronize
0.37%     15.719ms        8     1.9648ms  12.855us  5.7855ms  cudaMemcpy2DAsync
0.16%     6.9237ms        42    164.85us  9.2990us  1.2400ms  cudaMalloc
0.03%     1.3636ms        4     340.89us  336.01us  348.48us  cuDeviceTotalMem
0.02%     882.37us       352    2.5060us  243ns    71.116us  cuDeviceGetAttribute
0.01%     597.14us       114    5.2380us  623ns    159.46us  cudaEventCreateWithFlags
0.01%     543.28us       23    23.620us  10.228us  104.52us  cudaLaunch
0.01%     476.00us        6     79.332us  57.361us  124.03us  cudaMemcpy
0.01%     440.51us        4    110.13us  56.651us  165.33us  cudaStreamCreate
0.01%     350.52us        2    175.26us  56.939us  293.58us  cudaStreamCreateWithPriority
0.00%     103.51us        4     25.876us  15.561us  32.709us  cuDeviceGetName
0.00%     84.422us       32    2.6380us  624ns    7.1390us  cudaSetDevice
0.00%     71.908us      110     653ns    411ns    2.2620us  cuDeviceGetAttribute
0.00%     63.964us      147     435ns    256ns    1.6370us  cudaSetupArgument
0.00%     25.443us       23    1.1060us  496ns    3.7500us  cudaConfigureCall
0.00%     16.948us       10    1.6940us  1.1310us  2.3760us  cudaGetDevice
0.00%     11.324us        1    11.324us  11.324us  11.324us  cudaBindTexture
0.00%     9.1350us       16     570ns    365ns    938ns    cudaPeekAtLastError
0.00%     5.5870us        1     5.5870us  5.5870us  5.5870us  cudaStreamGetPriority
0.00%     5.0210us        6     836ns    275ns    2.1010us  cuDeviceGetCount
0.00%     4.0940us        2     2.0470us  1.4600us  2.6340us  cudaEventRecord
0.00%     3.6890us        2     1.8440us  1.4830us  2.2060us  cudaStreamWaitEvent
0.00%     3.5460us        2     1.7730us  1.6860us  1.8600us  cudaDeviceGetStreamPriorityRange
0.00%     3.5240us        6     587ns    283ns    1.1230us  cuDeviceGet
0.00%     3.3930us        3    1.1310us  1.0670us  1.1640us  cuInit
0.00%     3.0760us        6     512ns    289ns    760ns    cudaGetLastError
0.00%     2.1040us        3     701ns    650ns    744ns    cuDriverGetVersion
0.00%     1.5270us        1     1.5270us  1.5270us  1.5270us  cudaUnbindTexture
0.00%     1.2840us        1     1.2840us  1.2840us  1.2840us  cudaGetDeviceCount
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-8c359489-3b52-4933-8995-449377d0ac26.tar.gz. The data will be pre
sent for only a short duration of time.
* Server has ended your request.

```

Figure 2 screenshot of API calls.

Table 2 some time-consuming API calls

Time(%)	Time	Name
46.79%	1.9858s	cudaStreamCreateWithFlag
29.09%	1.2344s	cudaFree
20.47%	868.93ms	cudaMemGetInfo
2.99%	127.03ms	cudaStreamSynchronize

2.1 Simple CPU implementation

In this step, we implemented a CPU convolutional kernel. We flowed the forward convolution described in Chapter 16 of the textbook. The implementation is a for-loop that loop through all the computation positions and produces the convolution result. The classification results with our convolution kernel are presented in the following figures.

```

--0.11.0)
Installing collected packages: mxnet
  Running setup.py develop for mxnet
Successfully installed mxnet
* Running python /src/m2.1.py
Loading fashion-mnist data... done
Loading model... done
Op Time: 18.190609
Correctness: 0.8562 Model: ece408-high
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-3dbe502a-67f9-4044-ad5e-3a2d78393a0c.tar.gz. The data will be present for only a short duration of time.
* Server has ended your request.

```

Figure 3 ece408-high model execution time and accuracy

```
Installing collected packages: mxnet
  Running setup.py develop for mxnet
Successfully installed mxnet
* Running python /src/m2.1.py ece408-low 100
Loading fashion-mnist data... done
Loading model... done
Op Time: 0.201315
Correctness: 0.63 Model: ece408-low
* The build folder has been uploaded to http://s3.amazonaws.com/files.rai-project.com/userdata/build-7d417298-b51b-4a98-9005-7097a0b9ed36.tar.gz. The data will be present for only a short duration of time.
* Server has ended your request.
```

Figure 4 ece408-low model execution time and accuracy

Contribution: All the members thoroughly discussed the problem and distributed the task reasonably. Guanchen He(ghe10) wrote the CPU code; Yingyi Zhang(yingyiz2) performed the experiments; Guxin Jin(gjin7) wrote the report.